CSCIE-175 Final Project

Project: AWS
Name: Krunal Patel
E-Mail: k5patel@me.com

EC2 AMI (not avail after 1/13/2013): ami-85d35bec Github: https://github.com/lateralpunk/rgrt

YouTube: http://www.youtube.com/watch?v=9r9xSPlbVDA Website: http://lateralpunk.blogspot.com/2013/01/cscie-175-final-

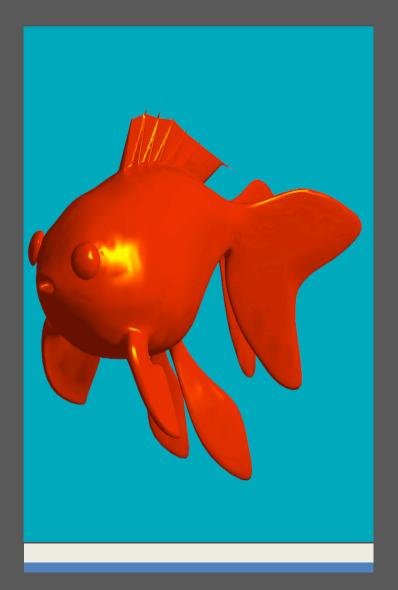
project-aws-cluster-gpu.html

Website of EC2 Instance (not avail after 1/13/2013): http://ec2-23-

20-133-181.compute-1.amazonaws.com/rgrt/

Website RGRT:

http://www.cs264.org/2009/projects/web/Patel_Krunal/index.html



AWS
CLUSTER
GPU RGRT

Introduction

I've based the last several years of my career focused in on GPU computing. It started off working in CUDA, and now I'm focused in on heterogeneous type computing using OpenCL. Back when I started, one of the major obstacles was obtaining access to compatible computers that can employ GPU computing. Hence I was delighted to learn about Amazon Web Service's offering of their Cluster GPU Instances.

The purpose of this project is to employ AWS Cluster GPU to implement a Regular-Grid Accelerated Ray Tracer. This RGRT will then be hosted on the same GPU instance to provide a web-service for rendering images. To make the project feasible, I decided employ the RGRT CUDA application I had worked on before. This would streamline my application development and also have something practical (and beautiful) to demo. One of the key things this was missing was sort of a service that end-users can employ to try out the ray tracer. I hope to achieve that goal by using AWS, Apache, and PHP.

Regular-Grid Accelerated Ray Tracing

The following explanation is derived from my previous work on this topic (http://www.cs264.org/2009/projects/web/Patel_Krunal/index.html).

Ray-tracing is a method used to draw realistic images on a computer. One can find examples of ray-traced images when viewing Disney/Pixar movies like Toy Story. The underlying algorithm behind ray-tracing is to decompose a picture into pixels, send rays through the pixels, intersect the rays with the objects in the scene, launch any secondary rays, and emit a final color for the original pixel.

Ray-tracing is known to be an embarrassingly parallel problem. Each pixel can be completely calculated independently of any other neighboring pixels. So a mechanism that allows spawning of multiple threads to parallelize the ray-tracing tasks would prove quite useful. Hence we have employed CUDA based technology to exemplify the performance gains to be had with GPU hardware acceleration for ray-tracing.

Our ray-tracer supports spheres, rectangles, and triangles as the basic primitives. Multiple light sources as well as shadows are available. Anti-aliasing via regular sampling is available too. Phong materials are utilized with support for glossy specular highlights. All world objects are composed into build scenes. Support for 3D model meshes by utilizing the PLY file format is available too.

Finally, though we get substantial improvements in rendering time taking a brute approach to ray-tracing, we also employ an acceleration scheme to the whole process. Obviously the brute approach is to traverse the whole scenegraph per pixel and the ray-tracing steps to find any intersections with world objects. A better option does exist. Our regular grid acceleration scheme subdivides the world into axis-aligned cells. These cells contain the world primitive objects that are within it's bounding box. The grid creation is done on the CPU. GPU grid traversal involve shooting a ray and intersecting with the cells and primitives that come in it's path. Once such an intersection is found, it is no longer necessary to continue with finding another intersection as we are guaranteed that the one just found is the closest. Thereafter shading proceeds as normal. More information about regular grid acceleration scheme can be found in the book "Ray Tracing from the Ground Up".

For the purposes of this project, we will treat the CUDA implementation of the RGRT as a black-box. We will assume it just works, and we are more interested in the plumbing involved in getting this to work on AWS as a web-service.

AWS Cluster GPU Setup

We wish to employ AWS to use the GPU horsepower. As a note, I would use Chrome browser for testing this project. To this extent, we list steps below to show our course of action. We will then go into each step in more detail:

- 1) Launch a Cluster GPU Instance that is not a spot instance
- 2) Setup the GPU instance with the right set of CUDA software packages
- 3) Setup the GPU instance with the right set of web hosting software packages
- 4) Get the source code for the RGRT off Github, build it and test it
- 5) Ensure that the PHP code implementing the web-service is up & running
- 6) Create an AMI of the above for safe-keeping (and so others can get at it quicker)
- 7) Publicly make the AMI available
- 8) Destroy the instance from 1)
- 9) Launch a Cluster GPU Spot Instance using the newly created AMI

So let the fun begin!!

First we need to launch the Cluster GPU instance that we will use as our setup machine to create the AMI. We decide not to employ a spot instance here because we don't want to have our instance taken from under our noses unexpectedly.

From our investigations, we know we want to employ the AMI: ami-02f54a6b:

Cluster Compute Amazon Linux AMI 2012.09

amazon/amzn-ami-gpu-hvm-2012.09.0.x86 64-ebs

hvm



achine

The Amazon Linux AMI 2012.09 is an EBS-backed, HVM image. It includes Linux 3.2, AWS tools, and repository access to multiple versions of MySQL, PostgreSQL, Python, Ruby, and Tomcat.

Root Device Size: 8 GB

lpdome.~ \$ ec2-describe-images -a | grep "ami-02f54a6b"



We then launch the instance from CLI:

ami-02f54a6b

lpdome.~ \$ ec2-run-instances	ami-02f54a6b -k k	krunalkp -t	cg1.4xlarge			
RESERVATION r-38808b40	739641513028	default				
INSTANCE i-33e87942	ami-02f54a6b		pending krunalkp	0	cg1.4xlarge	2
013-01-09T21:53:57+0000 us-e	ast-1b		monitoring-disabled			е
bs	hvm	xen	sg-5cc83b35 default	false		

Notice we need to explicitly set the Instance Type (-t cg1.4xlarge).

ebs

lpdome.~ \$ ec2-describe-instances i-33e87942 RESERVATION r-38808b40 739641513028 default INSTANCE rnal running krunalkp 0 i-33e87942 ami-02f54a6b ec2-107-21-135-16.compute-1.amazonaws.com ip-10-16-7-196.ec2.inte 2013-01-09T21:53:57+0000 cg1.4xlarge us-east-1b 10.16.7.196 ebs hvm sg-5cc83b35 default false BLOCKDEVICE /dev/sda1 vol-d96adda7 2013-01-09T21:54:01.000Z true

⋖	Name 🤏	Instance	AMI ID	Root Device	Туре	State	Status Checks	Alarm Status	Monitoring	Security Gr
⋖	empty	i-33e87942	ami-02f54a6b	ebs	cg1.4xlarge	running	Loading	none	basic	default



Now we ssh into the machine:

We do some updates first (sudo yum update) and reboot the machine. We install the following applications: emacs, git, freeglut-devel (for CUDA samples). I don't show the steps here for simplicity. We reboot (sudo reboot) after all the yum installs.

Next, by default the AMI is using CUDA 4.2. I wish to upgrade to the latest CUDA 5 since my code uses some utility functions only available in the later version of CUDA. But as a sanity check, let's ensure that we actually have CUDA compatible GPUs:

```
[ec2-user@ip-10-16-7-196 ~]$ lspci | grep -i nvidia
00:03.0 3D controller: NVIDIA Corporation GF100 [Tesla S2050] (rev a3)
00:04.0 3D controller: NVIDIA Corporation GF100 [Tesla S2050] (rev a3)
```

So as promised by AWS, we do have CUDA capable devices (2 for that matter). Next we verify that we have a compatible Linux distro (which we should ofcourse):

```
[ec2-user@ip-10-16-7-196 ~]$ uname -m && cat /etc/*release x86_64
Amazon Linux AMI release 2012.09
```

Ok, well this doesn't tell us much unfortunately. Doing a simple Google search gives us http://aws.amazon.com/amazon-linux-ami/2012.09-release-notes/

CUDA toolkit 4.2.9

The CUDA toolkit version 4.2.9 is available on the Cluster GPU AMI.

Ok well this proves that we don't have the right toolkit (and we will update), but still doesn't tell us the distro. This link http://ivan.manida.com/2012/11/installing-ffmpeg-on-amazon-ami-centos-6.html gives us the information we want: the AMI is a rebranded CentOS.

Ok then I need to verify gcc installed on the system:

```
[ec2-user@ip-10-16-7-196 ~]$ gcc --version
gcc (GCC) 4.6.2 20111027 (Red Hat 4.6.2-2)
Copyright (C) 2011 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

Oh nice, here we find out that the distro is Red Hat 4.6.2-2. We do have a good enough version of GCC as well. Finally, we go to the Nvidia CUDA page and download version 5:

LINUX: CUDA 5.0 Production Release									1
Getting Started					Guide R	elease Notes			
	Fedora	RI	HEL	Ubu	intu	OpenSUSE	SUSE	SUSE	
	16	5.X	6.X	11.10	10.04	12.1	Server 11 SP1	Server 11 SP2	
	64bit	64bit	64bit	64bit	64bit	64bit	64bit	64bit	
	32bit	32bit		32bit	32bit		32bit	32bit	J

Namely, we download RHEL 6.x 64-bit: http://developer.download.nvidia.com/compute/cuda/5_0/rel-update-1/installers/cuda_5.0.35_linux_64_rhel6.x-1.run

And then we install it:

```
[ec2-user@ip-10-16-7-196 tmp]$ sudo sh cuda_5.0.35_linux_64_rhel6.x-1.run Logging to /tmp/cuda install 1849.log
```

```
Do you accept the previously read EULA? (accept/decline/quit): accept
Install NVIDIA Accelerated Graphics Driver for Linux-x86_64 304.54? ((y)es/(n)o/(q)uit): y
Install the CUDA 5.0 Toolkit? ((y)es/(n)o/(q)uit): y
Enter Toolkit Location [ default is /usr/local/cuda-5.0 ]:
Install the CUDA 5.0 Samples? ((y)es/(n)o/(q)uit): y
Enter CUDA Samples Location [ default is /usr/local/cuda-5.0/samples ]:
```

Problem, there was a failure in installing the drivers. Let's check the error log:

ERROR: Unable to find the kernel source tree for the currently running kernel. Please make sure you have installed the kernel source files for your kernel and that they are properly configured; on Red Hat Linux systems, for example, be sure you have the 'kernel-source' or 'kernel-devel' RPM installed. If you know the correct kernel source files are installed, you may specify the kernel source path with the '--kernel-source-path' command line option.

Ok, so let's yum install kernel-devel:

```
[ec2-user@ip-10-16-7-196 tmp]$ sudo yum install kernel-devel
Loaded plugins: priorities, security, update-motd, upgrade-helper
Setting up Install Process
Resolving Dependencies
--> Running transaction check
 --> Package kernel-devel.x86 64 0:3.2.34-55.46.amzn1 will be installed
--> Finished Dependency Resolution
Dependencies Resolved
Package
                           Arch
                                                  Version
                                                                                    Repository
                                                                                                                 Size
Installing:
kernel-devel
                                                 3.2.34-55.46.amzn1
                            x86_64
                                                                                    amzn-updates
                                                                                                                7.3 M
Transaction Summary
Install 1 Package(s)
Total download size: 7.3 M
Installed size: 26_M
```

And now we retry CUDA 5 installer:

```
= Summarv =
Driver: Installed
Toolkit: Installed in /usr/local/cuda-5.0
Samples: Installed in /usr/local/cuda-5.0/samples (pristine) and /root/NVIDIA_CUDA-5.0_Samples (writable)
* Please make sure your PATH includes /usr/local/cuda-5.0/bin
Please make sure your LD LIBRARY PATH
   for 32-bit Linux distributions includes /usr/local/cuda-5.0/lib
   for 64-bit Linux distributions includes /usr/local/cuda-5.0/lib64:/lib
* OR
   for 32-bit Linux distributions add /usr/local/cuda-5.0/lib
   for 64-bit Linux distributions add /usr/local/cuda-5.0/lib64 and /lib
* to /etc/ld.so.conf and run ldconfig as root
* To uninstall CUDA, remove the CUDA files in /usr/local/cuda-5.0
* Installation Complete
Please see CUDA_Getting_Started_Guide_For_Linux.pdf in /usr/local/cuda-5.0/doc/pdf for detailed information on setting
up CUDA.
```

AWS CLUSTER GPU RGR1

page 8

Good, we got it. Let's update our .bash_profile file and then we reboot.

```
PATH=/usr/local/cuda-5.0/bin:$PATH:$HOME/bin
export PATH

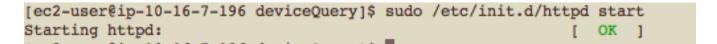
Export LD_LIBRARY_PATH=/usr/local/cuda-5.0/lib:/usr/local/cuda-5.0/lib64:$LD_LIBRARY_PATH
```

After reboot, we verify that CUDA is working by using the deviceQuery application (after running sudo make in the samples directory):

```
CUDA Device Query (Runtime API) version (CUDART static linking)
Detected 2 CUDA Capable device(s)
Device 0: "Tesla M2050"
 CUDA Driver Version / Runtime Version
                                                5.0 / 5.0
 CUDA Capability Major/Minor version number:
                                                 2.0
                                                2687 MBytes (2817982464 bytes)
 Total amount of global memory:
  (14) Multiprocessors x ( 32) CUDA Cores/MP:
                                                448 CUDA Cores
 GPU Clock rate:
                                                 1147 MHz (1.15 GHz)
 Memory Clock rate:
                                                 1546 Mhz
 Memory Bus Width:
                                                 384-bit
 L2 Cache Size:
                                                 786432 bytes
 Max Texture Dimension Size (x,y,z)
                                                 1D=(65536), 2D=(65536,65535), 3D=(2048,2048,2048)
 Max Layered Texture Size (dim) x layers
                                                 1D=(16384) x 2048, 2D=(16384,16384) x 2048
 Total amount of constant memory:
                                                 65536 bytes
 Total amount of shared memory per block:
                                                 49152 bytes
 Total number of registers available per block: 32768
 Warp size:
 Maximum number of threads per multiprocessor: 1536
 Maximum number of threads per block:
                                                 1024
 Maximum sizes of each dimension of a block:
                                                 1024 x 1024 x 64
 Maximum sizes of each dimension of a grid:
                                                 65535 x 65535 x 65535
 Maximum memory pitch:
                                                 2147483647 bytes
 Texture alignment:
                                                 512 bytes
 Concurrent copy and kernel execution:
                                                Yes with 2 copy engine(s)
 Run time limit on kernels:
                                                No
 Integrated GPU sharing Host Memory:
 Support host page-locked memory mapping:
 Alignment requirement for Surfaces:
                                                 Yes
 Device has ECC support:
                                                 Enabled
 Device supports Unified Addressing (UVA):
 Device PCI Bus ID / PCI location ID:
 Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >
Device 1: "Tesla M2050"
 CUDA Driver Version / Runtime Version
                                                5.0 / 5.0
 CUDA Capability Major/Minor version number:
                                                 2.0
 Total amount of global memory:
                                                 2687 MBytes (2817982464 bytes)
  (14) Multiprocessors x ( 32) CUDA Cores/MP:
                                                 448 CUDA Cores
 GPU Clock rate:
                                                1147 MHz (1.15 GHz)
 Memory Clock rate:
                                                 1546 Mhz
 Memory Bus Width:
                                                 384-bit
 L2 Cache Size:
                                                 786432 bytes
                                                 1D=(65536), 2D=(65536,65535), 3D=(2048,2048,2048)
 Max Texture Dimension Size (x,y,z)
 Max Layered Texture Size (dim) x layers
                                                1D=(16384) x 2048, 2D=(16384,16384) x 2048
 Total amount of constant memory:
                                                 65536 bytes
 Total amount of shared memory per block:
                                                 49152 bytes
 Total number of registers available per block: 32768
 Warp size:
 Maximum number of threads per multiprocessor: 1536
 Maximum number of threads per block:
                                                 1024
 Maximum sizes of each dimension of a block:
                                                 1024 x 1024 x 64
 Maximum sizes of each dimension of a grid:
                                                 65535 x 65535 x 65535
 Maximum memory pitch:
                                                 2147483647 bytes
 Texture alignment:
                                                 512 bytes
 Concurrent copy and kernel execution:
                                                Yes with 2 copy engine(s)
 Run time limit on kernels:
                                                 No
 Integrated GPU sharing Host Memory:
 Support host page-locked memory mapping:
                                                 Yes
 Alignment requirement for Surfaces:
                                                 Yes
 Device has ECC support:
                                                 Enabled
 Device supports Unified Addressing (UVA):
 Device PCI Bus ID / PCI location ID:
                                                 0 / 4
 Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >
deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 5.0, CUDA Runtime Version = 5.0, NumDevs = 2, Device0 = Tesla
M2050, Device1 = Tesla M2050
```

```
[ec2-user@ip-10-16-7-196 deviceQuery]$ sudo yum install httpd24 php54
Loaded plugins: priorities, security, update-motd, upgrade-helper
                                                                                                             00:00
amzn-gpu
                                                                                                 2.1 kB
amzn-main
                                                                                                 2.1 kB
                                                                                                             00:00
Setting up Install Process
Resolving Dependencies
--> Running transaction check
---> Package httpd24.x86 64 0:2.4.3-10.35.amzn1 will be installed
--> Processing Dependency: httpd24-tools = 2.4.3-10.35.amzn1 for package: httpd24-2.4.3-10.35.amzn1.x86_64
--> Processing Dependency: libaprutil-1.so.0()(64bit) for package: httpd24-2.4.3-10.35.amzn1.x86_64
--> Processing Dependency: libapr-1.so.0()(64bit) for package: httpd24-2.4.3-10.35.amzn1.x86_64
---> Package php54.x86_64 0:5.4.9-1.28.amzn1 will be installed
--> Processing Dependency: php54-common(x86-64) = 5.4.9-1.28.amzn1 for package: php54-5.4.9-1.28.amzn1.x86_64
--> Processing Dependency: php54-cli(x86-64) = 5.4.9-1.28.amzn1 for package: php54-5.4.9-1.28.amzn1.x86_64
--> Running transaction check
---> Package apr.x86 64 0:1.4.6-1.10.amzn1 will be installed
---> Package apr-util.x86_64 0:1.4.1-4.13.amzn1 will be installed
---> Package httpd24-tools.x86_64 0:2.4.3-10.35.amzn1 will be installed
---> Package php54-cli.x86_64 0:5.4.9-1.28.amzn1 will be installed
---> Package php54-common.x86 64 0:5.4.9-1.28.amzn1 will be installed
--> Finished Dependency Resolution
Dependencies Resolved
 Package
                             Arch
                                                    Version
                                                                                     Repository
Installing:
                                                    2.4.3-10.35.amzn1
 httpd24
                             x86_64
                                                                                     amzn-updates
                                                    5.4.9-1.28.amzn1
                                                                                     amzn-updates
                                                                                                                 3.1 M
 php54
                              x86 64
Installing for dependencies:
                                                                                                                 110 k
                              x86_64
                                                    1.4.6-1.10.amzn1
                                                                                    amzn-main
 apr-util
                                                    1.4.1-4.13.amzn1
                                                                                                                  87 k
 httpd24-tools
                             x86 64
                                                    2.4.3-10.35.amzn1
                                                                                    amzn-updates
                                                                                                                  83 k
 php54-cli
                                                    5.4.9-1.28.amzn1
                                                                                     amzn-updates
                                                                                                                 2.9 M
                              x86 64
 php54-common
                              x86_64
                                                    5.4.9-1.28.amzn1
                                                                                     amzn-updates
                                                                                                                 969 k
Transaction Summary
Install
             7 Package(s)
Total download size: 8.3 M
Installed size: 26 M
Is this ok [y/N]:
```

Then we manually startup the httpd server:



And check that Apache is running through our web-browser:



This page is used to test the proper operation of the Apache HTTP server after it has been installed. If you can read this page, it means that the web server installed at this site is working properly, but has not yet been configured.

If you are a member of the general public:

The fact that you are seeing this page indicates that the website you just visited is either experiencing problems, or is undergoing routine maintenance.

If you would like to let the administrators of this website know that you've seen this page instead of the page you expected, you should send them e-mail. In general, mail sent to the name "webmaster" and directed to the website's domain should reach the appropriate person.

For example, if you experienced problems while visiting www.example.com, you should send e-mail to "webmaster@example.com".

The <u>Amazon Linux AMI</u> is a supported and maintained Linux image provided by <u>Amazon Web Services</u> for use on <u>Amazon Elastic Compute Cloud (Amazon EC2)</u>. It is designed to provide a stable, secure, and high performance execution environment for applications running on <u>Amazon EC2</u>. It also includes packages that enable easy integration with <u>AWS</u>, including launch configuration tools and many popular AWS libraries and tools. <u>Amazon Web Services</u> provides ongoing security and maintenance updates to all instances running the <u>Amazon Linux AMI</u>. <u>The Amazon Linux AMI</u> is provided at no additional charge to <u>Amazon EC2 users</u>.

If you are the website administrator:

You may now add content to the directory <code>/var/www/html/</code>. Note that until you do so, people visiting your website will see this page, and not your content. To prevent this page from ever being used, follow the instructions in the file <code>/etc/httpd/conf.d/welcome.conf.</code>

You are free to use the images below on Apache and Amazon Linux AMI powered HTTP servers. Thanks for using Apache and the Amazon Linux AMI!

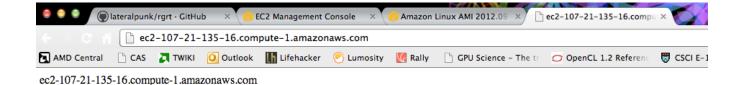


We want to ensure that the httpd service starts up on boot:

```
[ec2-user@ip-10-16-7-196 ~]$ sudo chkconfig httpd on
[ec2-user@ip-10-16-7-196 ~]$ chkconfig --list
acpid
                0:off
                        1:off
                                 2:on
                                         3:on
                                                 4:on
                                                         5:on
                                                                  6:off
atd
                0:off
                        1:off
                                 2:off
                                         3:on
                                                 4:on
                                                         5:on
                                                                  6:off
auditd
                0:off
                        1:off
                                 2:on
                                         3:on
                                                 4:on
                                                         5:on
                                                                  6:off
cloud-init
                0:off
                        1:off
                                 2:on
                                         3:on
                                                 4:on
                                                         5:on
                                                                  6:off
cloud-init-user-scripts 0:off
                                1:off
                                         2:on
                                                 3:on
                                                         4:on
                                                                 5:on
                                                                          6:off
crond
                0:off
                        1:off
                                 2:on
                                         3:on
                                                 4:on
                                                         5:on
                                                                  6:off
                0:off
httpd
                        1:off
                                 2:on
                                         3:on
                                                 4:on
                                                         5:on
                                                                 6:off
                       1:off
                                                                 6:off
ip6tables
                0:off
                                 2:on
                                        3:on
                                                 4:on
                                                         5:on
iptables
                0:off
                        1:off
                                 2:on
                                         3:on
                                                 4:on
                                                         5:on
                                                                 6:off
irqbalance
                0:off
                        1:off
                                 2:off
                                         3:on
                                                 4:on
                                                         5:on
                                                                 6:off
lvm2-monitor
                0:off 1:on
                                 2:on
                                        3:on
                                                 4:on
                                                         5:on
                                                                  6:off
                                                                  6:off
mdmonitor
                0:off
                       1:off
                                 2:on
                                         3:on
                                                 4:on
                                                         5:on
messagebus
                0:off
                        1:off
                                 2:on
                                         3:on
                                                 4:on
                                                         5:on
                                                                  6:off
                0:off
                                 2:off 3:off
                        1:off
                                                 4:off
                                                         5:off
                                                                  6:off
netconsole
netfs
                0:off
                        1:off
                                 2:off
                                         3:on
                                                 4:on
                                                         5:on
                                                                  6:off
network
                0:off
                        1:off
                                 2:on
                                         3:on
                                                 4:on
                                                         5:on
                                                                  6:off
ntpd
                0:off
                        1:off
                                 2:on
                                         3:on
                                                 4:on
                                                         5:on
                                                                 6 toff
                0:off
                       1:off
                                                 4:on
                                                                  6:off
ntpdate
                                2:on
                                        3:on
                                                         5:on
                                                         5:on
                                                                  6:off
nvidia
                0:off
                        1:off
                                 2:off
                                         3:on
                                                 4:on
psacct
                0:off
                        1:off
                                 2:off
                                         3:off
                                                 4:off
                                                         5:off
                                                                  6:off
racoon
                0:off
                        1:off
                                2:off 3:off 4:off
                                                         5:off
                                                                  6:off
                0:off
                        1:off
                                 2:off
                                         3:off
                                                 4:off
                                                         5:off
                                                                  6:off
rdisc
                0:off
                        1:off
                                 2:on
                                         3:on
                                                 4:on
                                                                  6:off
rsyslog
                                                         5:on
                0:off
                        1:off
                                                                  6:off
saslauthd
                                 2:off
                                         3:off
                                                 4:off
                                                         5:off
sendmail
                0:off
                        1:off
                                 2:on
                                         3:on
                                                 4:on
                                                         5:on
                                                                  6:off
sshd
                0:off
                        1:off
                                 2:on
                                         3:on
                                                 4:on
                                                         5:on
                                                                  6:off
udev-post
                0:off
                        1:on
                                 2:on
                                         3:on
                                                 4:on
                                                         5:on
                                                                  6:off
```

After reboot:

Now, let's remove the error page, and have just a single index.php file that will host our logic:



Ok, now let's go onto getting the source code for the RGRT from github where I posted it publicly. Here we use the https protocol so that we can provide our username & password. If we instead used the SSH protocol that the AMI image that I would eventually make would have my credentials!

```
[ec2-user@ip-10-16-7-196 ~]$ git clone https://github.com/lateralpunk/rgrt.git Cloning into rgrt...
remote: Counting objects: 72, done.
remote: Compressing objects: 100% (64/64), done.
remote: Total 72 (delta 7), reused 69 (delta 7)
Unpacking objects: 100% (72/72), done.
```

```
[ec2-user@ip-10-16-7-196 rgrt]$ pwd
/home/ec2-user/rgrt
[ec2-user@ip-10-16-7-196 rgrt]$ ls
build cuda html Makefile models README.md
[ec2-user@ip-10-16-7-196 rgrt]$
```

Let's ensure that we can build our application:

```
[ec2-user@ip-10-16-7-196 rgrt]$ pwd
/home/ec2-user/rgrt
[ec2-user@ip-10-16-7-196 rgrt]$ make
make[1]: Entering directory `/home/ec2-user/rgrt/cuda'
/usr/local/cuda/bin/nvcc -m64 -gencode arch=compute_10,code=sm_10 -gencode arch=compute_20,code=sm_20 -gencode arch=compute_30,code=sm_30 -gencode arch=compute_35,code=sm_35 -I/usr/local/cuda/include -I. -I. -I/usr/local/cuda/samples/common/inc -o main.o -c main.cu

g++ -m64 -o rgrt-cuda main.o bmploader.o plyfile.o -L/usr/local/cuda/lib64 -lcudart
mkdir -p ../bin/linux/release
cp rgrt-cuda ../bin/linux/release
make[1]: Leaving directory `/home/ec2-user/rgrt/cuda'
Finished building RGRT
```

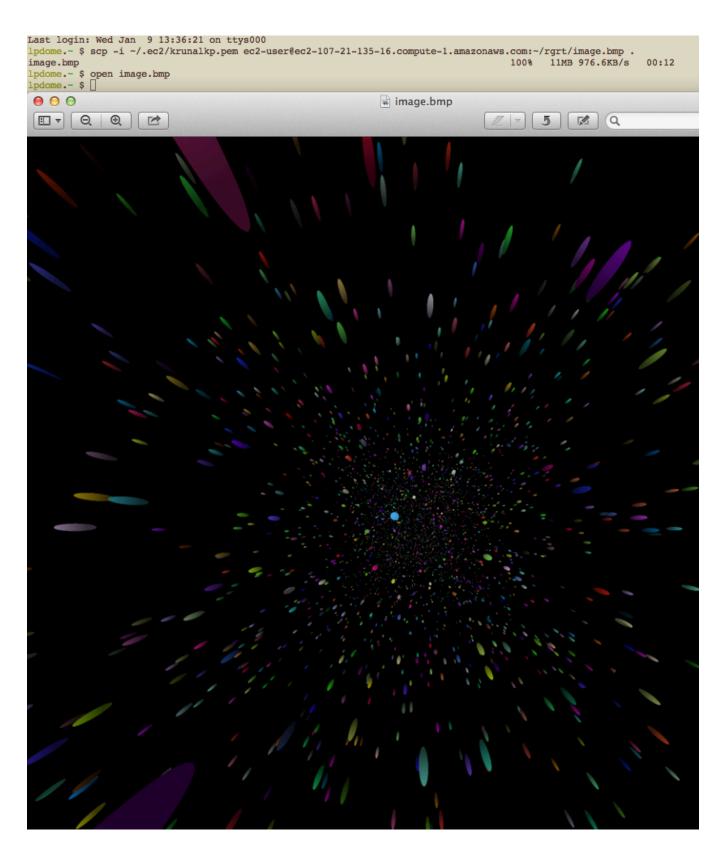
Ok, good, now launch it:

AWS CLUSTER GPU RGR1

page 14

```
[ec2-user@ip-10-16-7-196 rgrt]$ bin/linux/release/rgrt-cuda -b=build02 -g=1 -w=2000 -h=2000 -n=1 -o=image.bmp
*** Configuration - b: build02, o: image.bmp, g: 1, w: 2000, h: 2000, n: 1, z: 1.00, m: 0, s: 1, p: 0 ****
Starting ray-tracing...
Reorganizing Data Structures...
  Not going to use host pinned memory...
 Using Grid acceleration...
   Organizing grid cells...
     Stats: # of cells that have 0 objects, 1 object, 2 objects, etc...
     num_cells = 5832
     numZeroes = 302, numOnes = 759, numTwos = 1115
      numThrees = 1232 numGreater = 2424
   Finished grid cell organization.
Finished Consolidation.
World Stats:
 Total # of GeometricObjects: 24379
 Approximate bytes of data: 3901144 bytes
Starting CPU ray tracing...
Starting GPU ray tracing ...
 # of threads in a block: 16 x 8 (128)
 # of blocks in a grid : 125 x 250 (31250)
Calculating accuracy...
 Error: 0.008568
Imaged saved: image.bmp
Finished ray-tracing...
[ec2-user@ip-10-16-7-196 rgrt]$ ls image.bmp
```

Ok, well it look like it ran just fine, but we need to verify the output image.bmp, let's scp it to look at it:



Cool, looks like all the plumbing is in effect.

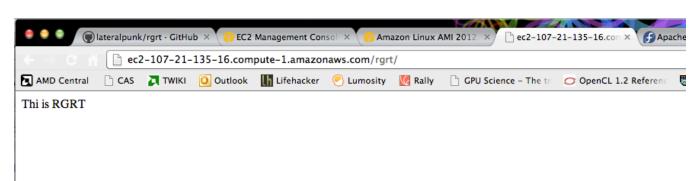
Now 1 more area of concern. Since we are employing Apache, we need to make a alias to a directory that is web server accessible to our html pages. Instead of copying the files over to /var/www/html, I instead decided to use mod_alias. The first thing I did was copy the rgrt git repo over to the root directory. Then I ensured that /rgrt/www is world accessible:

```
[ec2-user@ip-10-16-7-196 www]$ pwd
/rgrt/www
[ec2-user@ip-10-16-7-196 www]$ ls -la
total 12
drwx--x--x 2 ec2-user ec2-user 4096 Jan 9 23:39 .
drwx--x--x 8 ec2-user ec2-user 4096 Jan 10 00:37 ..
-rwxr-xr-x 1 ec2-user ec2-user 11 Jan 9 23:39 index.php
[ec2-user@ip-10-16-7-196 www]$
```

Then we update the httpd.conf file:

```
<IfModule alias module>
   # Redirect: Allows you to tell clients about documents that used to
   # exist in your server's namespace, but do not anymore. The client
   # will make a new request for the document at its new location.
   # Redirect permanent /foo http://www.example.com/bar
   # Alias: Maps web paths into filesystem paths and is used to
   # access content that does not live under the DocumentRoot.
   # Example:
   # Alias /webpath /full/filesystem/path
   # If you include a trailing / on /webpath then the server will
   # require it to be present in the URL. You will also likely
   # need to provide a <Directory> section to allow access to
   # the filesystem path.
   # ScriptAlias: This controls which directories contain server scripts.
   # ScriptAliases are essentially the same as Aliases, except that
   # documents in the target directory are treated as applications and
   # run by the server when requested rather than as documents sent to the
   # client. The same rules about trailing "/" apply to ScriptAlias
   # directives as to Alias.
   ScriptAlias /cgi-bin/ "/var/www/cgi-bin/"
   Alias /rgrt /rgrt/www
</IfModule>
```

Then we can finally access our website (there's nothing there right now):



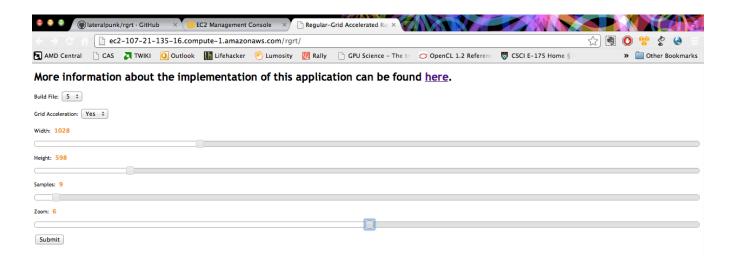
Now, we implement the web service in PHP. This service will feature a form that allows a user to select a build scene and adjust some paramaters:

```
//ensure that the script doesn't time out
set_time_limit(0);
if (isset($_POST['doit'])) {
 //defaults
$build = $_POST['build'];
 $g = $_POST['grid'];
$width = $_POST['width'];
 $height = $_POST['height'];
$n = $_POST['samples'];
 $Z = $_POST['zoom'];
$0 = "/rgnt/www/images/" . uniqid('rgrt-', true) . "-rgnt.bmp";
  //ok execute the rart-cuda pro
 shell_exec("export LD_LIBRARY_PATH-/usr/local/cuda-5.0/lib:/usr/local/cuda-5.0/lib64:\$LD_LIBRARY_PATH;cd /rgrt;bin/linux/release/rgrt-cuda m=1 -g=". $g ." -b=" . $build .
 -w=" . $width . " -h=" . $height . " -n=" . $n . " -z=" . $z . " -p=0 -o=" . $o);
 //now use the Linux at command to delete the newly created file after 15 minutes shell_exec("export SHELL=/bin/bash && echo \underline{m} -f " . $o . " | at now + 15 minutes");
 sleep(2);
 //now show the user the image:
header('Location: images' . strrchr($0,'/'));
<!doctype html>
<html lang="en">
<head>
 <meta charset="utf-8" />
 <title>Regular-Grid Accelerated Ray Tracing</title>
link rel="stylesheet" href="http://code.jquery.com/ui/1.9.2/themes/base/jquery-ui.css" />
 <script src="http://code.jquery.com/jquery-1.8.3.js"></script>
<script src="http://code.jquery.com/ui/1.9.2/jquery-ui.js"></script>

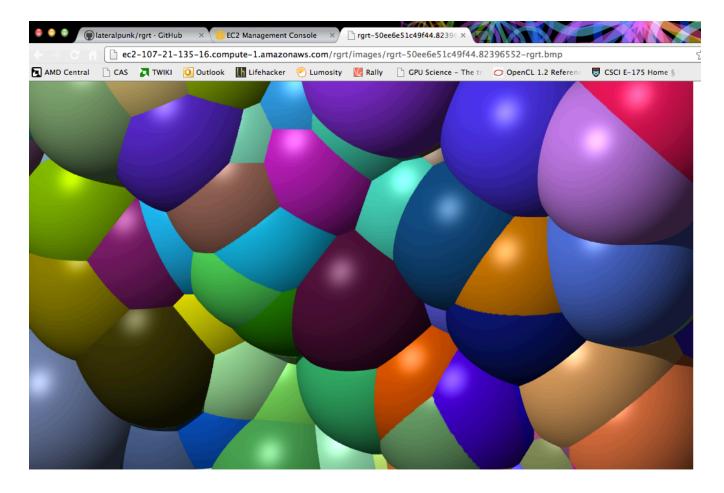
 <script>
 $(function() {
    $( "#slider-width" ).slider({
     range: "max",
         min: 16.
          value: 512,
         step: 2,
slide: function( event, ui ) {
$( "#width" ).val( ui.value );
```

```
});
     $( "#width" ).val( $( "#slider-width" ).slider( "value" ) );
   });
 $(function() {
     $( "#slider-height" ).slider({
   range: "max",
       min: 16,
       max: 4096,
       value: 512,
       step: 2,
       slide: function( event, ui ) {
       $( "#height" ).val( ui.value );
     }
   });
     $( "#height" ).val( $( "#slider-height" ).slider( "value" ) );
   });
 $(function() {
     $( "#slider-samples" ).slider({
   range: "max",
       min: 1,
       max: 256,
       value: 1,
       step: 1,
       slide: function( event, ui ) {
       $( "#samples" ).val( ui.value );
   });
     $( "#samples" ).val( $( "#slider-samples" ).slider( "value" ) );
   });
 $(function() {
     $( "#slider-zoom" ).slider({
   range: "max",
       min: -1024,
       max: 1024,
       value: 1,
       step: 1,
       slide: function( event, ui ) {
       $( "#zoom" ).val( ui.value );
     }
   });
     $( "#zoom" ).val( $( "#slider-zoom" ).slider( "value" ) );
   });
 </script>
</head>
<body>
<form action="index.php" method="post">
<input type="hidden" name="doit" value="doit">
```

```
<label>Build File:</label>
<select name="build">
<option value="build01">1</option>
<option value="build02">2</option>
<option value="build03">3</option>
<option value="build04">4</option>
<option value="build05">5</option>
<option value="build06">6</option>
<option value="build07">7</option>
<option value="build08">8</option>
<option value="build09">9</option>
</select>
<label>Grid Acceleration:</label>
<select name="grid">
<option value="1">Yes</option>
<option value="0">No</option>
</select>
<label>Width:</label>
 <input type="text" id="width" name="width" style="border: 0; color: #f6931f; font-weight: bold;" />
<div id="slider-width"></div>
 <label>Height:</label>
 <input type="text" id="height" name="height" style="border: 0; color: #f6931f; font-weight: bold;" />
<div id="slider-height"></div>
 <label>Samples:</label>
 <input type="text" id="samples" name="samples" style="border: 0; color: #f6931f; font-weight: bold;" />
<div id="slider-samples"></div>
 <label>Zoom:</label>
 <input type="text" id="zoom" name="zoom" style="border: 0; color: #f6931f; font-weight: bold;" />
<div id="slider-zoom"></div>
<input type="submit" value="Submit">
</form>
</body>
</html>
```



And after a little while:



After 15 minutes the above file will automatically delete itself off the server. Try different paramaters and you'll see that everything is dynamic.

Ok, now that the application is complete, let's create an AMI out of this:

```
lpdome.~ $ ec2-create-image -n rgrt_image i-33e87942
IMAGE ami_85d35bec
```

We see that the ssh that we were connected to before gets disconnected:

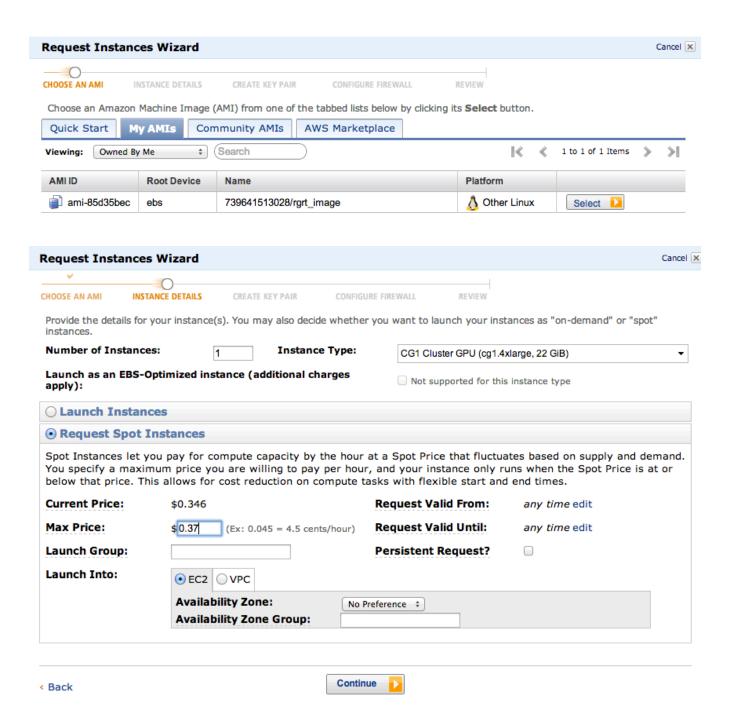
Checking AWS EC2 we get the following:

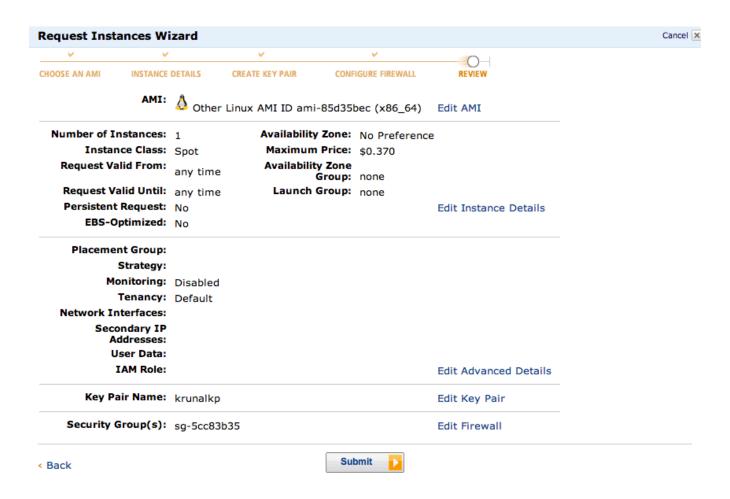


After waiting for a little while, our AMI has been created:



Ok good. Now let's launch a spot instance of the AMI we just created. We are going to leave this 1 running so that you can evaluate it. We'll shut down the original instance so that we don't get charged crazy. **NOTE: change of plan, I** wanted to document all this down right now, but I'll start up an instance that is running right before I submit my assignment to you so that I don't incur unnecessary charges.





And then we wait for a spot to appear:

Request ID	Max Price	AMI ID	Instance	Туре	State	Status
sir-02429812	\$0.370	ami-85d35bec		cg1.4xlarge	open	pending-evaluation

Got it:

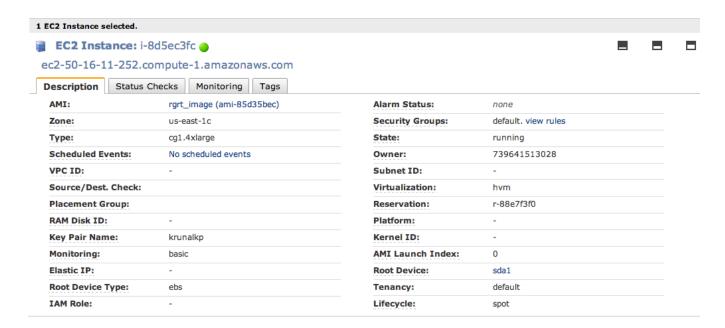
AWS CLUSTER GPU RGR1

page 25

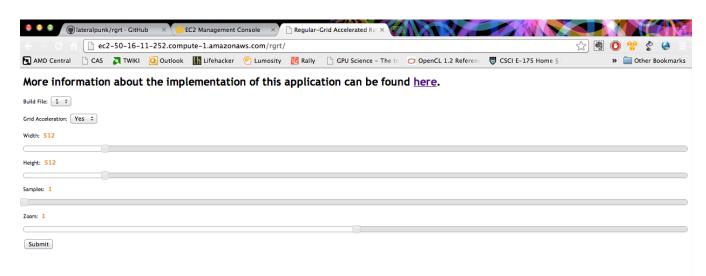
	Request ID	Max Price	AMI ID	Instance	Туре	State	Status
⋖	sir-02429812	\$0.370	ami-85d35bec	i-8d5ec3fc	cg1.4xlarge	active	fulfilled

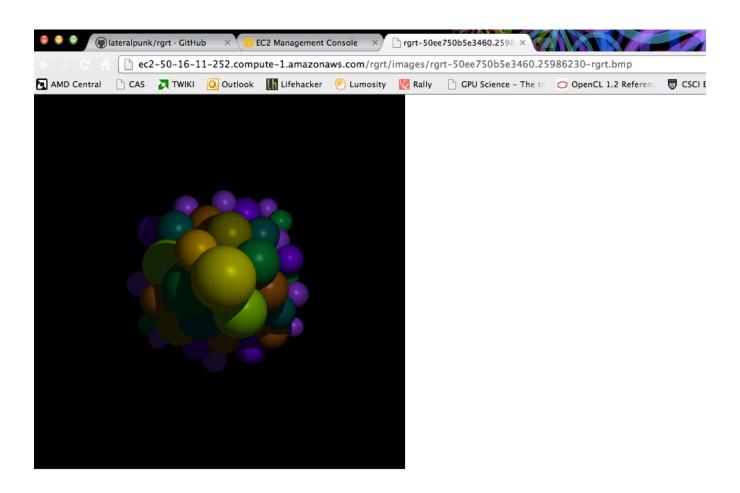
1 Spot Instance Request selected AMI: ami-85d35bec Request Persistence: one-time Instance Type: Availability Zone: cg1.4xlarge **Monitoring Enabled:** false Availability Zone Group: Maximum Price: \$0.370 Launch Group: Request Valid From: Request Valid Until: **Key Pair Name:** krunalkp Security Group(s): sg-5cc83b35 Status Code: fulfilled Status Update Time: 2013-01-09 23:55 PST Status Message: Your Spot request is fulfilled. Created: 2013-01-09 23:52 PST State: active State Reason: Instance: i-8d5ec3fc Subnet: **Product Description:** Linux/UNIX Kernel ID: RAM Disk ID: IAM Role: Launched availability zoneus-east-1c





At this point, it should be a simple matter of point our browser to the new URL:





And everything just magically worked! Awesome. Now let's terminate our original instance since it's costing us too much:

Terminate Instances

Are you sure you want to terminate this instance?

Warning: On an EBS-backed instance, the default action is for the root EBS volume to be deleted when the instance is terminated. Storage on any local drives will be lost.

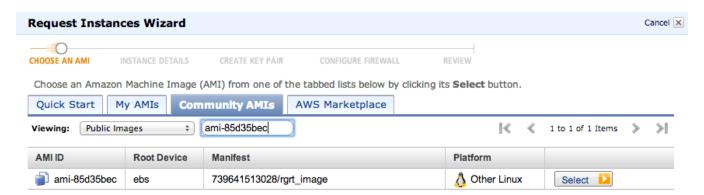
i-33e87942 (ec2-107-21-135-16.compute-1.amazonaws.com)

Cancel Yes, Terminate

Finally, so that you can evaluate the AMI, we make it's visibility Public:



And you can confirm it is public by:



So you can just use the above (make sure instance type is cg1.4xlarge) and everything should work just nicely.

Conclusion

Hence we've shown complete steps on how to setup an AWS EC2 GPU Instance and support our RGRT.